

Learning Objectives

- Use CSP to optimize resource usage for job-shop scheduling
- Use CSP to implement a scheduler for the TerraBot
- Implement a dynamic scheduler that is based on agent observations

This is a group assignment consisting of three programming parts, plus a presentation. You should decide with your teammates how to divide the work on the first three parts (you can have one person do each part, or you can work on all the parts together), but everyone **is to work together** on Part 4 (the presentation).

Download and unzip the assignment code from Canvas. Place the folder in `~/Desktop/TerraBot/agents/`. You will also need a library called `dill`, which is built on top of `pickle`. Install it using `pip install dill`.

In this assignment, you will make extensive use of `ortools`. Some of the common functions you may want to use to add constraints include:

- `Add()`: used to indicate (in)equality between variables (or constants), multiplying/dividing by constants, summing variables in a list (remember that you can sum binary variables to get counts).

You can add `.OnlyEnforceIf()` to the end of an `Add` constraint to indicate that the constraint should be applied only if the argument to `OnlyEnforceIf` is `True` (or 1). Note that the argument **must be a simple Boolean variable**, for instance, it cannot be a sum or a relation. If you need to, you can get that behavior by creating an intermediate Boolean variable and setting it equal to whatever formula you need to be true. For instance:

```
model.Add(foo == sum(z_vars))
model.Add(x == y).OnlyEnforceIf(foo)
```

- `AddMultiplicationEquality()`: used to constrain one variable equal to the product of other variables (`AddDivisionEquality` works similarly).
- `AddMaxEquality()`: constrain one variable to the maximum of a list of variables (`MinEquality` works similarly).
- `AddNoOverlap()`: constrain a set of **interval variables** not to overlap in time (that is, for each interval in the list, either the start of an interval is \geq the ends of all the other intervals, or the end of that interval is \leq the starts of the other intervals).

When using these constraints with `OptionalIntervals`, `ortools` does the “right” thing by enforcing these constraints only for intervals that are “present” (i.e., those that are scheduled to occur).

- `AddReservoirConstraintWithActive()`: Used to model resource usage. The function takes a list of time variables, at which points some change to the resource might occur; a list of “demands”, which are the changes to the resources; a list of “actives”, which are Boolean variables that indicate whether the change is actually scheduled to take place, and a min/max range of the resource.

The idea is that, for each time, if the associated active variable (same position in the list) is `True/1`, then the associated demand is added to/subtracted from the reservoir (depending on whether the demand is positive or negative). The `AddReservoirConstraintWithActive` constraint ensures that the value of the reservoir never exceeds the range limits (i.e., is never less than the minimum and never greater than the maximum). When you set up the reservoir, you don’t need to worry about the order of the lists (e.g., which time comes first) - the constraint will take care of that automatically for you. See the November 11 lecture notes for more details.

Refer to https://developers.google.com/optimization/reference/python/sat/python/cp_model for a complete list of the functions you can use to add constraints.

Part 1: Resource Utilization for Automated Factories (27 points)

In this part, you will be modifying `job_scheduler.py` by using `ortools` to schedule factory jobs while trying to maximize profit. A factory job produces an object (or objects) with a certain value, and one maximizes profit by using various resources in optimal ways to produce high-value objects. Specifically, a **job** consists of a list of tasks that must be completed in the order given. A **task** can be carried out using a **machine** and, for most tasks, there are options about which machine to use (note that different tasks might need to use the same machine). The decision about which machine to use for a task is based on the availability of the machine, the resources that it uses, the time it takes the machine to complete the task, and the value of the product produced (which relates to how much money the object can be sold for). Typically, there is an inverse relationship between time to complete the task and value – the longer it takes to complete the task, the higher the quality, and hence the higher the value of the object produced.

A task typically requires a set of **tools** and a set of **parts** to complete the task. Tools are in use during the task, but then get returned to the pool; Parts are used up by the task, but there may be jobs that can produce additional parts. A task also defines which machines can be used to complete the task.

There is also a **deadline** for completing the various tasks (e.g., an 8-hour shift at the factory). It is often the case that not all requested jobs can be completed within the allotted time – part of the scheduling problem is to determine which jobs to schedule and which to leave out.

We have defined a number of classes in `job_scheduler.py` to aid in this assignment:

- **Machine:** represents a factory machine that can be used to make things. A **Machine** object has a **name** and its **energy_cost** per hour (some machines are more efficient than others).
- **Tool:** represents a pool of tools that can be used by tasks, such as some number of wrenches. A **Tool** object has a **name** and the number of tools (**num**) available in the pool.
- **Part:** represents a pool of parts that are needed to produce objects. A **Part** object has a **name**, the **quantity** of parts initially available in the pool, and the **cost** of using one of the parts.
- **Task:** represents part of the sequence needed by a job to make objects. A **Task** object has a **name**, a list of **tools** that it requires, and a list of **parts** it requires. Each tool in the list needs to be available throughout the duration of the task. Each part in the list needs to be available at the start of the task and is used up by the task. If a task needs multiple tools or parts of the same type, they are repeated in the list (e.g. ['wrench', 'hammer', 'wrench'] indicates the need for two wrenches and a hammer). The **task_machines** attribute is a list of **TaskMachine** instances (see below) that indicate which machines can achieve the task.
- **PartsTask:** represents a task that can create new parts, which can then be used by other tasks that need that type of part. It is a subtype of **Task**. In addition to the **name**, **tools** and **parts** required for a **Task**, a **PartsTask** object specifies which **part** is produced and the **quantity** that is produced at one time (i.e., executing the task once may produce multiple parts).
- **TaskMachine:** represents the fact that a machine can be used to complete some task. A **TaskMachine** specifies the **Task** and **Machine** objects, the **duration** the machine takes to complete the task (given by an integral number of hours), and the **value** of the product produced.
- **Job:** Represents the tasks needed to create some object. A **Job** specifies the **name** of the job and a list of **Tasks**, all of which need to be done, in order, to produce the object. Note that the value of a job is equal to the sum of the values for all the tasks that comprise that job.

For example, one might have job J1 consisting of tasks T1 and T2. T1 needs tool **Tool1** and parts **Part1** and **Part2**, and can be done by either machine M1, taking one hour and producing a value of 200, or by M2, taking 2 hours and producing a value of 300. Task T2 needs no tools and two instances of part **Part3** and can be done by either machine M2, taking 3 hours and producing a value of 500, or by M3, taking 2 hours and producing a value of 250. Thus, there are 4 choices for completing J1: 1) using M1 and M2, taking 4 hours with a value of 700; 2) using M1 and M3, taking 3 hours with a value of 450; 3) using M2 followed by M2 again, taking 5 hours with a value of 800; and 4) using M2 and M3, taking 4 hours with a value of 550. Which is best

depends on what other jobs need to be scheduled and how much energy each machine uses (since electricity usage reduces profit).

We have also specified sets of variables that will likely be useful in what follows (see the function `create_job_task_variables` in `job_scheduler.py`). Specifically:

- **starts:** is a dictionary whose keys are a list of job, task and machine. The values are the integer `ortools` variables representing the start times for doing the task on the particular machine for that job.
- **ends:** is a dictionary whose keys are a list of job, task and machine, and whose values are the integer variables representing the end times for doing the task on the particular machine for that job.
- **scheduleds:** is a dictionary whose key a list of job, task and machine, and whose values are Boolean `ortools` variables indicating whether that machine was actually scheduled to do the task for that job.
- **intervals:** is a dictionary whose keys are (job, task, machine) tuples, and whose values are `OptionalInterval` variables that encode the start, end, and duration of a task being performed by the machine. All intervals are constrained to start at, or after, time 1 and end no later than the deadline.

Note: for consistency, use the function `self._key(Job, Task, Machine)` to generate keys to access the above dictionaries.

The following steps create your constraint model one step at a time. At each step, you can test your model using `python autograder.py -p 1 -s <n>`, where `n` is the step number. All constraints up to, and including, that step will be added to test your solution. You can also use the `-v` flag to generate more verbose output and the `-g` flag to visualize the schedule (although for some of the steps, where the constraints are not fully developed, visualization may not work correctly).

Step 1: Task Constraints (2 points)

The first step is to implement `create_task_constraints` to encode the constraint that, for each job, each task must be achieved by only one machine. Recall that the definition of tasks include a list of `TaskMachine` objects that indicate what machines can be used to achieve that task.

Test using `python autograder.py -p 1 -s 1`

Note: If you use the `-v` option, it will print out your solution; If you use the `-g` option, it will display your solution graphically (you can use these options for all the subsequent steps, as well).

Step 2: Machine Constraints (2 points)

The next step is to implement `create_machine_constraints` to encode the constraint that each machine can handle only one task at a time. The implementation will be a bit different from the class scheduling exercise you did in class since here we are using the interval model, rather than the binary model, but the lecture slides should help.

Test using `python autograder.py -p 1 -s 2`, which will also include the task constraints from Step 1.

Step 3: Task Ordering Constraints (3 points)

This constraint is a bit more involved. You need to implement `create_task_ordering_constraints` to encode that all the tasks of a given job are done in sequence – for instance, whichever machine is used for the first task in a job, it must end before the machine chosen for the second task in that job begins. Recall that the `scheduleds` variables indicate whether a given machine has actually been scheduled to achieve a task for a job.

Note that the list of `task_machines` is a `TaskMachine` instance (in past years, many students initially thought they were `Machine` instances). Also, this may be a good place to use `OnlyEnforceIf`.

Test using `python autograder.py -p 1 -s 3`

Step 4: Task Completion Constraints (3 points)

This constraint is also a bit more involved. You need to implement `create_task_completion_constraints` to encode that if any job is started it must be finished by the deadline. That is, either *all* tasks in a job are scheduled (more accurately, the machines scheduled to achieve the tasks), or *none* of them are.

Test using `python autograder.py -p 1 -s 4`

Step 5: Tool Constraints (5 points)

Now we get to resource usage! For this step, implement `create_tool_constraints` to encode how tools are used by the factory. That is, at the start of a scheduled task the tool is removed from the pool and at the end it is returned. Use the `AddReservoirConstraintWithActive()` constraint to ensure that the number of tools in concurrent use is never greater than the pool size. See the lecture slides from November 11 if you need help on this part.

Test using `python autograder.py -p 1 -s 5`

Step 6: Part Constraints (5 points)

Continuing with resource usage, implement `create_part_constraints` to encode how parts are used by the factory. That is, if a scheduled task needs a part, it is removed from the pool of parts at the **start** of the task (and is not returned at the end – it is consumed). In addition, if a scheduled `PartsTask` task creates a part, the quantity of that part produced is added to the pool at the **end** of the `PartsTask` task (you can determine if a task is a `PartsTask` with the function `task.isPartsTask()`). Use the `AddReservoirConstraintWithActive()` constraint to ensure that the number of parts never falls below zero. **As a extra constraint, assume that the number of parts in the pool can never be greater than the number at the start** – this encodes that there is some storage limit for parts. See the lecture notes from November 11 if you need help on this part.

Test using `python autograder.py -p 1 -s 6`

Step 7: Maximizing Profit (5 points)

Up until now, the optimization criterion was to maximize total value (implemented for you in `add_values`). For this step, you will maximize net profit, which is defined as the **value** of the products minus the **cost** of producing them. There are two costs that need to be taken into account:

- Each machine uses energy to produce the parts (some machines are more efficient than others). One part of the costs is the total cost of the energy used by the machine over the duration that it is working (note that the `energy_cost` attribute of a `Machine` is the energy usage **per hour**).
- Tasks may need to use parts in the manufacturing process (the `Task` object specifies the list of parts it needs to use). Those parts have a cost, specified by `part.cost` that must be accounted for in the cost function.

Fill in the `add_costs` function to set the `ortools` variable `self.cost` to the total cost of production (the sum of the two costs above). The function `add_optimization` will then define an objective function that maximizes net profit (you don't need to alter `add_optimization`). **Hint:** if you multiply some equation by a `scheduleds` variable, the result will be zero if the variable is `False` (0) and the value of the equation if the variable is `True` (1).

Test using `python autograder.py -p 1 -s 7`

Step 8: Testing Everything Together (2 points)

When you are all done, run `python autograder.py -p 1`, and it will run all the tests for this part. In particular, it will test whether constraints you added in later steps somehow messed up earlier solutions. It

will also test a larger range of problems – some of which will take a significant amount of time to run. Don't forget that you can use the `-v` and `-g` options to get more insight into what your code is producing.

Part 2: Creating a Greenhouse Scheduler (28 points total)

In this part, you will modify `greenhouse_scheduler.py` by using `ortools` to create a program that will schedule each of your behaviors. We present below four steps that will guide you through a way to write the scheduler, but you are free to delete all our code in `greenhouse_scheduler.py` and create your own if you wish.

We have implemented the `GreenhouseScheduler` class that has the basic functions for this part, including creating the necessary constraint variables in the `__init__` function. In particular, `self.all_jobs[behavior, time]` is a Boolean constraint of whether a behavior is enabled at the time block time. `self.minutes_per_chunk` represent how many minutes are in a block of time and `self.horizon` represents how many chunks of time there are in a day. If not `None`, `self.sched_file` is the name of the file where the generated schedule will be saved. Finally, `self.behaviors_info` is a dictionary of all the behaviors. For each behavior (key), the value is a (t, s, m) tuple where:

- t:** The *least* cumulative amount of time, in minutes, that all scheduled behaviors of that type need to be run.
- s:** The spacing between the behaviors, in minutes, given by (min, max). Specifically, no two behaviors can be scheduled for closer than “min” minutes and at least one behavior must be scheduled in every block of “max” minutes.

Note that this constraint does not wrap around midnight (i.e., the constraint does not apply from the end of one day to the beginning of the next).

- m:** The maximum amount of time the behavior should run at night between hours [20,24) and [0,8), that is, starting from 8pm until just before 8am.

For instance, assuming that behaviors are enabled for 30 minute at a time (as specified by `self.minutes_per_chunk`), a value for `behavior_info['LowerHumid']` of “(480, (60, 120), 300)” means that the “LowerHumid” behavior needs to be run for at least 480 minutes (16 thirty-minute blocks) per day, instances must be separated by at least 60 minutes (2 blocks), at least one instance of the behavior must be scheduled every 120 minutes (4 blocks), and “LowerHumid” behaviors can be scheduled for at most 300 minutes (10 blocks of time) at night (8pm to 8am). Note that a “min” time of 0 means that behaviors can be scheduled back-to-back.

The auto-grader has 5 test cases, the first checks to see whether your model's constraints are feasible (that is, they produce *some* solution) and the second checks to see whether they produce no solution if the constraints are infeasible. The final three tests check whether you are producing a solution that is consistent with the `refsol`'s constraints.

Step 1: Implementing Duration Constraints (6 points)

For this step you will implement the duration constraints of each behavior (that is, how many instances are scheduled) by filling in the function `createDurationConstraints`.

For this, you should loop through each of the behaviors in `behavior_info` and make sure that the sum of all the times that behavior is scheduled to run is at least the time required by the first index of `behavior_info`. Make sure to convert the time in minutes specified by `behavior_info` into the number of required time blocks using `self.minutes_per_chunk`.

You can test this with `python autograder.py -p 2 -s 1`

Note: If you use the `-v` option, it will print out your solution; If you use the `-g` option, it will display your solution graphically (you can use these options for all the subsequent steps, as well).

Step 2: Mutually Exclusive Constraints (6 points)

For this step you will implement the mutually exclusive constraints between behaviors (that is, behaviors that need the same resource should not be scheduled at the same time), by filling in the function `createMutuallyExclusiveConstraints`.

For this, you should loop through each time block from 0 until the time horizon and make sure that none of the mutually exclusive behaviors are enabled at the same time. A behavior is mutually exclusive with another behavior if it shares the same actuator, if one behavior wants the actuator on and the other off, or if it is the Raise and Lower for the same environmental value (temperature, humidity, moisture). A full list is provided as a comment in the function. **Hint:** A simple way to do this is to simply include all the pairwise mutually exclusive constraints.

You can test everything so far with `python autograder.py -p 2 -s 2`

Step 3: Night Constraints (6 points)

For this step you will implement the night constraint of each behavior (that is the maximum amount of time the behavior can operate at night) by filling in the function `createNightConstraints`.

For this, you should loop through each of the behaviors in `behavior_info` and make sure that the sum of all the times of that behavior run during the night is less than or equal to the time allowed for that behavior by the third index of the `behavior_info` tuple. Make sure to convert the time in minutes specified by `behavior_info` into the number of required time blocks using `self.minutes_per_chunk`.

Note: “Night” is considered to run from 8pm to 8am; since this wraps at midnight, you need to combine the blocks from 8pm to midnight with those from midnight to 8am.

You can test everything so far with `python autograder.py -p 2 -s 3`

Step 4: Spacing Constraints (9 points)

For this step you will implement the min and max spacing constraints of each behavior by filling in the function `createSpacingConstraints`.

For this, you should loop through each of the behaviors and add constraints that ensure there is at least the minimum time between scheduled instances and that there is at least one instance scheduled in every maximum block of time.

A suggested approach is to first convert the min and max spacing values to blocks of time `min_b` and `max_b` (using `self.minutes_per_chunk`). Then, for all windows of `min_b+1` consecutive blocks, **at most** one instance of the behavior can be scheduled, and for all windows of `max_b+1` blocks **at least** one instances must be scheduled. If a behavior is not run at night, create min/max windows for that behavior only during the day (8am to 8pm). Make sure to convert the times in minutes to the desired number time blocks. Finally, note that the windows at the end of the day may be smaller than the min/max times (e.g., a window starting at 11pm can last for only an hour). You don’t have to worry about the max constraint for blocks of time that would extend to after either midnight or 8pm, depending on whether the behavior runs at night.

You can test everything with `python autograder.py -p 2 -s 4`

Step 5: Testing Everything Together (1 point)

When you are all done, run `python autograder.py -p 2`, and it will run all the tests for this part. In particular, it will test whether constraints you added in later steps somehow messed up earlier solutions. Don’t forget that you can use the `-v` and `-g` options to get more insight into what your code is producing.

Part 3: Dynamic Scheduler (30 points)

For this part, your team needs to come up with an approach to adapt the schedule to current conditions. This can take any number of forms – we offer two suggested approaches below. If you want to take it in another direction, please talk to us about your plan before Monday November 28, so that we can approve of what you intend to do.

Ideas:

Plants have different needs depending on their stage of growth, especially how much light and moisture they need/can tolerate. Do some research to determine what radishes and lettuce need when they first germinate, start growing, and get to the point where they are starting to mature (after about two weeks). Use vision to determine what stage your plants are at and have the schedule change the frequency that the lighting and watering behaviors are enabled. Also use the insolation prediction that you developed previously to monitor and maintain the appropriate amount of light during the day. Based on these values, you can adjust the next day's schedule such that the daily light and moisture average out appropriately.

You will notice that some behaviors rarely need to be enabled (e.g., lowering temperature rarely has to run, especially in the evening) and some behaviors may need to be run more frequently. Implement monitors to determine when behaviors are actually needed (e.g., they turn actuators on) and when they should be needed (e.g., some sensor limit is reached when the associated behavior is not running). Use this data, say averaged over a day, to change the scheduled frequency of behaviors. You don't have to do this for all the behaviors – pick two that you think would most benefit from this rescheduling.

To do the rescheduling, you can use the scheduler developed in Part 2 or, to help you, we have provided a pre-compiled scheduler that can be accessed using `greenhouse_scheduler_ref.py`. Read the introduction of Part 2 to see how to define the arguments to the `GreenhouseScheduler` class.

In all cases, you will need to:

In a file of your own choosing, create a `Monitor` to observe and compute the values you need, and have it update the `behavior_info` parameters needed for the scheduler. Given those updated parameters, call `GreenhouseScheduler` and `solveProblem` (a method of `GreenhouseScheduler`) to create a schedule based on the modified constraints, which in turn are based on observed conditions. You may find that you have to change the parameters in the example “main” schedule for some behaviors to allow you to optimize the behaviors you care about in this assignment. Call the Planner functions `setTestingSchedule` and `switch_to_test_sched` to set and load a testing file and run from that (rather than having it just read the txt file).

Part 4: Presentation (15 points)

For this part you will **work as a group** to create a presentation to present to the class the dynamic scheduler you have created and why you have chosen to do so and what benefits your scheduler has.

Run your agent in simulation using the standard `greenhouse_schedule.txt` for 2 (simulated) weeks. Then run your agent using your dynamic schedules.

Slide 1 and 2: Describe your goal in dynamic scheduling (idea 1, 2, or something else), and your implementation of it. What did your monitor do? What are you trying to dynamically optimize? How did you change the `GreenhouseScheduler` scheduler to capture this goal?

Slide 3 and 4: Capture 2-3 schedules and visualize them using the function `plot_binary.py` in `visualize_solution.py` (look at `autograder.py` for an example of how to use that function). Identify the constraints that made them different from the “main” schedule, and show visually how are they different from that “main” schedule.

Slides 5 and 6: Compare and contrast the plant growth based on these two techniques. Capture your estimated plant health or other quantities important for your scheduling. If possible, capture some visualizations of your plants in both cases to compare as well.

Submission

Submit the following files to Canvas: `greenhouse_scheduler.py`, `job_scheduler.py`, `README`, your presentation file (Powerpoint or PDF), and any files that you created/edited for Part 3.

When added to a directory with the rest of the agent files, your code should run without error.

The presentations will be on Monday, December 2. Everyone is expected to attend class that day and participate in your group's presentation.